

```
In [1]: # 1. Classification Using Hand-Crafted Features
# (a)
# Load VizWiz dataset
import os
import json
import requests
from pprint import PrettyPrinter

base_url = 'https://ivc.ischool.utexas.edu/VizWiz/data'
img_dir = '%s/Images/' % base_url
print(img_dir)

train_split = 'train'
train_file = '%s/Annotations/%s.json' % (base_url, train_split)
train_data = requests.get(train_file, allow_redirects=True)
print(train_file)

test_split = 'test'
test_file = '%s/Annotations/%s.json' % (base_url, test_split)
test_data = requests.get(test_file, allow_redirects=True)
print(test_file)

val_split = 'val'
val_file = '%s/Annotations/%s.json' % (base_url, val_split)
val_data = requests.get(val_file, allow_redirects=True)
print(val_file)

https://ivc.ischool.utexas.edu/VizWiz/data/Images/
https://ivc.ischool.utexas.edu/VizWiz/data/Annotations/train.json
https://ivc.ischool.utexas.edu/VizWiz/data/Annotations/test.json
https://ivc.ischool.utexas.edu/VizWiz/data/Annotations/val.json
```

```

In [2]: # Read the local file
training_data = train_data.json()
testing_data = test_data.json()
validation_data = val_data.json()
print("Length of training data:", len(training_data))
print("Length of test data:", len(testing_data))
print("Length of validation data:", len(validation_data))

image_name_train = []
question_train = []
label_train = []

image_name_val = []
question_val = []
label_val = []

image_name_test = []
question_test = []
label_test = []

num_train_VQs = 20000
for vq in training_data[0:num_train_VQs]:
    image_name_train.append(vq['image'])
    question_train.append(vq['question'])
    label_train.append(vq['answerable'])

num_val_VQs = 8000
for vq in validation_data[0:num_val_VQs]:
    image_name_val.append(vq['image'])
    question_val.append(vq['question'])
    label_val.append(vq['answerable'])

num_test_VQs = 3173
for vq in testing_data[0:num_test_VQs]:
    image_name_test.append(vq['image'])
    question_test.append(vq['question'])
#     label_test.append(vq['answerable'])

import pandas as pd
image_name_train = pd.DataFrame(image_name_train, columns=['image'])
image_name_val = pd.DataFrame(image_name_val, columns=['image'])
image_name_test = pd.DataFrame(image_name_test, columns=['image'])
question_train = pd.DataFrame(question_train, columns=['question'])
question_val = pd.DataFrame(question_val, columns=['question'])
question_test = pd.DataFrame(question_test, columns=['question'])

X_train = pd.concat([image_name_train, question_train], axis=1)
y_train = pd.DataFrame(label_train, columns=['label'])
X_val = pd.concat([image_name_val, question_val], axis=1)
y_val = pd.DataFrame(label_val, columns=['label'])
X_test = pd.concat([image_name_test, question_test], axis=1)
# y_test = pd.DataFrame(label_test, columns='label')

```

Length of training data: 20000

Length of test data: 8000

Length of validation data: 3173

```

In [12]: # (b)
# Use Microsoft Azure API to extract image-based features
subscription_key_vision = '412bc41b5b5844febf4d7cd63510fb4f'
vision_base_url = 'https://westcentralus.api.cognitive.microsoft.com/vision/v1.0'
vision_analyze_url = vision_base_url + '/analyze?'
from time import sleep

def analyze_image(image_url):
    # Microsoft API headers, params, etc
    headers = {'Ocp-Apim-Subscription-key': subscription_key_vision}
    params = {'visualfeatures': 'Description, Tags'}
    data = {'url': image_url}
    # send request, get API response
    try:
        response = requests.post(vision_analyze_url, headers = headers, params=params, json=data)
    except:
        sleep(10)
        response = requests.post(vision_analyze_url, headers = headers, params=params, json=data)
    # response = requests.post(vision_analyze_url, headers=headers, params=params, json=data)
    if (response.status_code == 200):
        analysis = response.json()
    else:
        print("get image {} failed".format(image_url))
        analysis = {"description": {"tags": []}}
    return analysis

def extract_features(data):
    return {
        'tags': data['description']['tags'],
        # 'confidence': data['tags'][0]['confidence']
    }

image_feature = {}
def get_image_feature(X):

    for i in range(20000):
        image_url = img_dir + '%s' % (X['image'][i])
        data = extract_features(analyze_image(image_url))
        tag_i = []
        for item in data['tags']:
            tag_i.append(item)
        tag_i_join = ' '.join(tag_i)
        # image_feature.append(tag_i_join)
        image_feature[str(i)] = tag_i_join
        if (i%500==0):
            print('get number', str(i))

    return image_feature
image_feature = get_image_feature(X_train)

```

```

get number 12500
get number 13000
get number 13500
get number 14000
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000014307.jpg failed
get number 14500
get number 15000
get number 15500
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000015541.jpg failed
get number 16000
get number 16500
get number 17000
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000017089.jpg failed
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000017311.jpg failed
get number 17500
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000017821.jpg failed
get number 18000
get number 18500
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000018603.jpg failed
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000018777.jpg failed
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000018938.jpg failed
get number 19000
get number 19500
get image https://ivc.ischool.utexas.edu/VizWiz/data/Images/VizWiz_train_000000019757.jpg failed

```

In [13]: *# Write image feature to csv file*

```

import csv

data = pd.DataFrame()
indexlist = []
featurelist = []
for index, feature in image_feature.items():
    indexlist.append(index)
    featurelist.append(feature)
data["id"] = indexlist
data["image_feature"] = featurelist
data.columns = ["id", "image_feature"]
data.head()
data.to_csv('image_feature_train.csv', index=False)

```

```

In [ ]: # Extract text features using Microsoft Azure
from time import sleep
subscription_key_text = 'e25225c679e74f61a2ab61924b41a866'
text_analytics_base_url = 'https://centralus.api.cognitive.microsoft.com/text/analytics/v2.0/'
key_phrase_api_url = text_analytics_base_url + 'keyPhrases'
question_feature = {}
def get_question_feature(question_train):

    for i in range(20000):

        question_json = question_train['question'][i]
        documents = {'documents': [{'id': i, 'text': question_json}]}
        headers = {"Ocp-Apim-Subscription-Key": subscription_key_text}
        maxiter = 10

        try:
            response = requests.post(key_phrase_api_url, headers = headers, json=documents)
        except:
            sleep(10)
            response = requests.post(key_phrase_api_url, headers = headers, json=documents)
        if(response.status_code == 200):
            question_json = response.json()['documents']
            question = pd.DataFrame(question_json)['keyPhrases']
            question = question.tolist()[0]
            tag_i=[]
            for item in question:
                tag_i.append(item)
            question = ' '.join(tag_i)
            question_feature[str(i)] = question
        else:
            print("not get",str(i))
            question_feature[str(i)] = ""
        if (i%500==0):
            print('get number',str(i))

    return question_feature
question_feature = get_question_feature(X_train)
#print(question_feature)

```

```

In [ ]: # Write key phrase to csv file
data = pd.DataFrame()
indexlist = []
keywordlist = []
for index,keyword in question_feature.items():
    indexlist.append(index)
    keywordlist.append(keyword)
data["id"] = indexlist
data["question_keyword"] = keywordlist
data.columns = ["id", "question_keyword"]
data.head()
data.to_csv('question_feature_train.csv', index=False)

```

In [ ]:

```
In [1]: # Load dataset
import pandas as pd

image_feature_train = pd.read_csv('dataset/image_feature_train.csv', header=None)[1].fillna(value='')
question_feature_train = pd.read_csv('dataset/question_feature_train.csv', header=None)[1].fillna(value='')
y_train = pd.read_csv('dataset/y_train.csv', header=None)[1].fillna(value='')

image_feature_val = pd.read_csv('dataset/image_feature_val.csv', header=None)[1].fillna(value='')
question_feature_val = pd.read_csv('dataset/question_feature_val.csv', header=None)[1].fillna(value='')
y_val = pd.read_csv('dataset/y_val.csv', header=None)[1].fillna(value='')

image_feature_test = pd.read_csv('dataset/image_feature_test.csv', header=None)[1].fillna(value='')
question_feature_test = pd.read_csv('dataset/question_feature_test.csv', header=None)[1].fillna(value='')

print(y_train.shape)
print(y_val.shape)

(20000,)
(2000,)
```

```
In [2]: # One hot encoding
import numpy as np
from sklearn.feature_extraction.text import CountVectorizer

def one_hot_transform(text):
    count = CountVectorizer()
    bag = count.fit_transform(text).toarray()
    return bag

image_feature = pd.concat([image_feature_train, image_feature_val, image_
_feature_test], axis=0)
question_feature = pd.concat([question_feature_train, question_feature_v
al, question_feature_test], axis=0)

image_feature_transformed = pd.DataFrame(one_hot_transform(image_feature
))
question_feature_transformed = pd.DataFrame(one_hot_transform(question_f
eature))

X = pd.concat([image_feature_transformed, question_feature_transformed],
axis=1)

X_train = X[:20000]
X_val = X[20000:22000]
X_test = X[22000:22100]

print(X_train.shape)
print(X_val.shape)
print(X_test.shape)

(20000, 3572)
(2000, 3572)
(100, 3572)
```



```
In [3]: # Dimension reduction
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
%matplotlib inline

# X_train = X_train[:4000]
# y_train = y_train[:4000]

principal_com = []
accuracy = []
for i in range(1, 1800, 60):
    pca = PCA(n_components=i)
    pca.fit(X_train)
    X_train_reduced = pca.transform(X_train)
    X_val_reduced = pca.transform(X_val)
    model = LogisticRegression()
    model.fit(X_train_reduced, y_train)
    y_val_predicted = model.predict(X_val_reduced)
    accuracy.append(accuracy_score(y_val, y_val_predicted))
    principal_com.append(i)

plt.plot(principal_com, accuracy)
plt.xlabel("Principal Components")
plt.ylabel("Accuracy")
plt.show()
```

[illegible]

[illegible]

2. Specify a solver to silence this warning.


FutureWarning)

/anaconda3/lib/python3.7/site-packages/sklearn/linear\_model/logistic.p

y:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2

2. Specify a solver to silence this warning.

FutureWarning)



```
In [5]: pca = PCA(n_components=61)
pca.fit(X_train)
X_train_reduced = pca.transform(X_train)
X_val_reduced = pca.transform(X_val)

model = LogisticRegression()
model.fit(X_train_reduced, y_train)
y_val_predicted = model.predict(X_val_reduced)

print("Accuracy on test set: {:.2f}".format(accuracy_score(y_val_predicted, y_val)))
```

Accuracy on test set: 0.71

/anaconda3/lib/python3.7/site-packages/sklearn/linear\_model/logistic.p

y:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2

2. Specify a solver to silence this warning.

FutureWarning)



(e)

I use the image-based features and question-based features as the input data. To extract these features, I use Microsoft Azure API (computer vision and text analytics) to process all the VizWiz training dataset, 2000 rows of validation dataset, and 100 rows of test dataset. I use all the image tags from Azure Vision API, connect them word by word and store it as image feature. Then I use Text Analytics API to extract key phrases of each question and store it as question-based feature.

After getting all necessary data, I create the following dataset for training and testing. There are three image feature datasets of training, validation and testing. Three question feature datasets of training, validation and testing. Two output datasets of training and validation. Then I combine image and question feature datasets together and use one-hot to encode. That's how I prepare all data following the following model training.

First I use PCA to reduce the dimension of input data. Then I train the neural network, logistic regression, bagging, boost, and SVM models using the training data. During this process, I tuned the model parameters to make sure each model works well and has a relatively accuracy on the validation data. After that, I choose the one with the highest accuracy (some of them has similar accuracy, I just randomly pick one among them) and make prediction of the test dataset using that model.

(f)

I tried neural network, logistic regression, bagging, AdaBoost, SVM models. I trained the model with 20000 rows of training data, then tested the model on 2000 rows of validation data. The accuracy of different models are as follows. MLPClassifier: 0.64; Logistic Regression: 0.71; bagging: 0.64; boost: 0.71; SVM: 0.71. According to the accuracy on validation data, I choose SVM to make prediction on the 100 examples from VizWiz\_test000000020100.jpg to VizWiz\_test\_000000020199.jpg.

Here are some hyperparameters of different models.

MLPClassifier: hidden\_layer\_sizes=(2048,4096,4096), max\_iter=1000, random\_state=42, activation='relu', solver='adam'

BaggingClassifier: max\_samples=50

AdaBoostClassifier: base\_estimator=DecisionTreeClassifier(max\_depth=10), n\_estimators=500, learning\_rate=0.1

Logistic regression, boost, and SVM models have same accuracy on validation dataset, we can apply any of them to the 100 examples of the test split.

```
In [2]: # Classification using neural networks
        # (a)
        import pandas as pd
        from tensorflow.examples.tutorials.mnist import input_data
        from sklearn.model_selection import train_test_split

        mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
        X = mnist.train.images
        y = mnist.train.labels.astype("int")
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25
        , random_state=42)
```

Extracting MNIST\_data/train-images-idx3-ubyte.gz

Extracting MNIST\_data/train-labels-idx1-ubyte.gz

Extracting MNIST\_data/t10k-images-idx3-ubyte.gz

Extracting MNIST\_data/t10k-labels-idx1-ubyte.gz

```
In [3]: # (b)
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier

# Standardization
stdsc = StandardScaler()
X_train_std = stdsc.fit_transform(X_train)
X_test_std = stdsc.fit_transform(X_test)

for num_layer in range(1, 9):

    for num_neuron in range(5, 13):
        mlp = MLPClassifier(hidden_layer_sizes = ((num_neuron,)*num_layer), max_iter=50, random_state=42, activation='relu', solver='adam', batch_size=200)
        mlp.fit(X_train_std, y_train)
        print("Number of hidden layers", num_layer)
        print("Number of Neurons per layer", num_neuron)
        print("Accuracy on test set: {:.5f}".format(mlp.score(X_test_std, y_test)))
    #         print(mlp.coefs_)
```



```
/home/nbuser/anaconda3_501/lib/python3.6/site-packages/sklearn/neural_network/multilayer_perceptron.py:564: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (50) reached and the optimization hasn't converged yet.  
  % self.max_iter, ConvergenceWarning)
```

Number of hidden layers 1  
Number of Neurons per layer 5  
Accuracy on test set: 0.74713  
Number of hidden layers 1  
Number of Neurons per layer 6  
Accuracy on test set: 0.77469  
Number of hidden layers 1  
Number of Neurons per layer 7  
Accuracy on test set: 0.81978  
Number of hidden layers 1  
Number of Neurons per layer 8  
Accuracy on test set: 0.84873  
Number of hidden layers 1  
Number of Neurons per layer 9  
Accuracy on test set: 0.86204  
Number of hidden layers 1  
Number of Neurons per layer 10  
Accuracy on test set: 0.87033  
Number of hidden layers 1  
Number of Neurons per layer 11  
Accuracy on test set: 0.88080  
Number of hidden layers 1  
Number of Neurons per layer 12  
Accuracy on test set: 0.88022  
Number of hidden layers 2  
Number of Neurons per layer 5  
Accuracy on test set: 0.79040  
Number of hidden layers 2  
Number of Neurons per layer 6  
Accuracy on test set: 0.82364  
Number of hidden layers 2  
Number of Neurons per layer 7  
Accuracy on test set: 0.86407  
Number of hidden layers 2  
Number of Neurons per layer 8  
Accuracy on test set: 0.87658  
Number of hidden layers 2  
Number of Neurons per layer 9  
Accuracy on test set: 0.88640  
Number of hidden layers 2  
Number of Neurons per layer 10  
Accuracy on test set: 0.89324  
Number of hidden layers 2  
Number of Neurons per layer 11  
Accuracy on test set: 0.90153  
Number of hidden layers 2  
Number of Neurons per layer 12  
Accuracy on test set: 0.90953  
Number of hidden layers 3  
Number of Neurons per layer 5  
Accuracy on test set: 0.79556  
Number of hidden layers 3  
Number of Neurons per layer 6  
Accuracy on test set: 0.84276  
Number of hidden layers 3  
Number of Neurons per layer 7  
Accuracy on test set: 0.87607

Number of hidden layers 3  
Number of Neurons per layer 8  
Accuracy on test set: 0.89375  
Number of hidden layers 3  
Number of Neurons per layer 9  
Accuracy on test set: 0.89716  
Number of hidden layers 3  
Number of Neurons per layer 10  
Accuracy on test set: 0.90524  
Number of hidden layers 3  
Number of Neurons per layer 11  
Accuracy on test set: 0.91135  
Number of hidden layers 3  
Number of Neurons per layer 12  
Accuracy on test set: 0.91098  
Number of hidden layers 4  
Number of Neurons per layer 5  
Accuracy on test set: 0.46305  
Number of hidden layers 4  
Number of Neurons per layer 6  
Accuracy on test set: 0.83949  
Number of hidden layers 4  
Number of Neurons per layer 7  
Accuracy on test set: 0.85687  
Number of hidden layers 4  
Number of Neurons per layer 8  
Accuracy on test set: 0.88582  
Number of hidden layers 4  
Number of Neurons per layer 9  
Accuracy on test set: 0.89113  
Number of hidden layers 4  
Number of Neurons per layer 10  
Accuracy on test set: 0.90458  
Number of hidden layers 4  
Number of Neurons per layer 11  
Accuracy on test set: 0.90873  
Number of hidden layers 4  
Number of Neurons per layer 12  
Accuracy on test set: 0.91433  
Number of hidden layers 5  
Number of Neurons per layer 5  
Accuracy on test set: 0.68393  
Number of hidden layers 5  
Number of Neurons per layer 6  
Accuracy on test set: 0.83404  
Number of hidden layers 5  
Number of Neurons per layer 7  
Accuracy on test set: 0.86175  
Number of hidden layers 5  
Number of Neurons per layer 8  
Accuracy on test set: 0.87738  
Number of hidden layers 5  
Number of Neurons per layer 9  
Accuracy on test set: 0.88902  
Number of hidden layers 5  
Number of Neurons per layer 10  
Accuracy on test set: 0.90924

Number of hidden layers 5  
Number of Neurons per layer 11  
Accuracy on test set: 0.91324  
Number of hidden layers 5  
Number of Neurons per layer 12  
Accuracy on test set: 0.91745  
Number of hidden layers 6  
Number of Neurons per layer 5  
Accuracy on test set: 0.81025  
Number of hidden layers 6  
Number of Neurons per layer 6  
Accuracy on test set: 0.65251  
Number of hidden layers 6  
Number of Neurons per layer 7  
Accuracy on test set: 0.84175  
Number of hidden layers 6  
Number of Neurons per layer 8  
Accuracy on test set: 0.88015  
Number of hidden layers 6  
Number of Neurons per layer 9  
Accuracy on test set: 0.89527  
Number of hidden layers 6  
Number of Neurons per layer 10  
Accuracy on test set: 0.90611  
Number of hidden layers 6  
Number of Neurons per layer 11  
Accuracy on test set: 0.91011  
Number of hidden layers 6  
Number of Neurons per layer 12  
Accuracy on test set: 0.90691  
Number of hidden layers 7  
Number of Neurons per layer 5  
Accuracy on test set: 0.54880  
Number of hidden layers 7  
Number of Neurons per layer 6  
Accuracy on test set: 0.83615  
Number of hidden layers 7  
Number of Neurons per layer 7  
Accuracy on test set: 0.84800  
Number of hidden layers 7  
Number of Neurons per layer 8  
Accuracy on test set: 0.87382  
Number of hidden layers 7  
Number of Neurons per layer 9  
Accuracy on test set: 0.85789  
Number of hidden layers 7  
Number of Neurons per layer 10  
Accuracy on test set: 0.90022  
Number of hidden layers 7  
Number of Neurons per layer 11  
Accuracy on test set: 0.90495  
Number of hidden layers 7  
Number of Neurons per layer 12  
Accuracy on test set: 0.90727  
Number of hidden layers 8  
Number of Neurons per layer 5  
Accuracy on test set: 0.58625

```
Number of hidden layers 8
Number of Neurons per layer 6
Accuracy on test set: 0.76356
Number of hidden layers 8
Number of Neurons per layer 7
Accuracy on test set: 0.86007
Number of hidden layers 8
Number of Neurons per layer 8
Accuracy on test set: 0.87753
Number of hidden layers 8
Number of Neurons per layer 9
Accuracy on test set: 0.89789
Number of hidden layers 8
Number of Neurons per layer 10
Accuracy on test set: 0.89069
Number of hidden layers 8
Number of Neurons per layer 11
Accuracy on test set: 0.89811
Number of hidden layers 8
Number of Neurons per layer 12
Accuracy on test set: 0.90422
```

```
In [9]: mlp = MLPClassifier(hidden_layer_sizes = (12, 12, 12, 12, 12), max_iter=
50, random_state=42, activation='relu', solver='adam', batch_size=200)
mlp.fit(X_train_std, y_train)
weights = mlp.coefs_
print(len(weights))
```

6

```
In [10]: for i in range(6):
print(weights[i].shape)
```

```
(784, 12)
(12, 12)
(12, 12)
(12, 12)
(12, 12)
(12, 12)
(12, 10)
```

```
In [11]: print(784*12+12*12*4+12*10)
```

10104

(c)

The optimal hyperparameter: number of hidden layer = 5, number of neurons per layer = 12

The number of weights is  $784 \times 12 + 12 \times 12 \times 4 + 12 \times 10 = 10104$

(d)

The performance of the neural network largely depends on the number of hidden layers and number of neurons per layer. We can improve the performance of the neural network by increasing the number of hidden layers and number of neurons per layer. If the number is too small, the model could be underfitting.

When the number of hidden layer and neurons per layer reach certain amount, the performance of the neural network won't improve too much if we keep enlarging that number. The appropriate number depends on the complexity of our training data. If the number is too big, it might cause the accuracy to decrease since it could be overfitting.